

LA-5: More Data Wrangling (15 points)

Learning Outcomes

In this assignment, you will learn how to:

- Run and examine frequency distributions using the `freq()` function in the `summarytools` package.
- View descriptive statistics using the `descr()` function in the `summarytools` package.
- Filter for a range of values using `>=`, `<=`, and `&`.
- Create new variables from existing variables in a data frame using the `mutate()` function.
- Recode variables using the `case_when()` function.

Tip

Read all the instructions carefully before starting the assignment. Set up your R script as you did in LA-4 (name, uNID, date, assignment number).

Important

How to report your answers. Write each answer as a comment in your R script, labeled with its question number (e.g., # 2a: ...). Round percentages and statistics to two decimal places, and include units of measurement (e.g., kW, meters) wherever a variable has them.

Instructions

- 1) Install (`install.packages()`) and load (`library()`) the packages below. Then download the `wind-turbines.csv` dataset and its codebook from Canvas, upload the dataset to your Posit Cloud project (as you did in LA-4), and read it into R.

- `tidyverse`
- `summarytools`
- `rstatix`

Tip

Just like in previous assignments, read the data into an object using `read_csv()`. For example: `wind <- read_csv("wind-turbines.csv")`. Take a few minutes to look through the codebook so you know what each variable measures before you start answering questions below.

- 2) The dataset includes data about wind turbines in the US. **Each row (case) is one wind turbine.** Use the `freq()` function to examine the frequency distributions of the variables `Site.State` and `Year`. Read the note below the questions before you start, then use the frequency distributions to answer the following questions.

- a) How many wind turbines are located in Utah? (States are coded with two-letter abbreviations; Utah is UT.)
- b) How many wind turbines became operational in 2021?
- c) What percentage of turbines became operational between 1983 and 1999 (include these years in the calculation)?

Note

You can use `freq()` with the pipe, just like the functions you used in LA-4: `wind |> freq(variable)`. (You may also see it written as `freq(wind$variable)`; both do the same thing.)

The output is a table listing each category (or, for a numeric variable like `Year`, each value). The parts you will use are:

- The **Freq** column: the count of cases (turbines) with that value.
- The **% Valid** column: the percentage of cases with that value. (This dataset has no missing values, so **% Valid** and **% Total** are the same.)
- The **Total** row at the bottom: the total number of cases in whatever data you ran `freq()` on.

To answer (c), you'll need to combine two skills:

- **Filtering on a range of values.** So far, you've filtered using `==` to match one exact value (e.g., `Show.Name == "Mamma Mia!"`). You can also use `>=` ("greater than or equal to") and `<=` ("less than or equal to") to filter for a range, and combine two conditions in one `filter()` step with `&` ("and"). For example, to keep only cases where `Year` is 1983 or later **and** 1999 or earlier:

```
wind |>
  filter(variable >= value1 & variable <= value2) |>
  freq(variable)
```

- **Calculating a percentage.** Once you filter for `Year` between 1983 and 1999 and run `freq(Year)` on that filtered data, look at the **Total** row to see how many turbines became operational in that range. Then run `freq(Year)` on the full (unfiltered) dataset and look at its **Total** row to see the total number of turbines overall. Divide the first number by the second, and multiply by 100, to get a percentage.

- 3) Now, subset the data to answer the following question: How many turbines in Utah have been operational since 2010?

Tip

This requires combining two conditions with `&`, the same technique described above: one condition for the state ("UT") and one for the year. "Since 2010" means 2010 and every year after, so use `>=`. (Using two separate `filter()` steps, as in LA-4, also works.)

Save the filtered data as a new object so you can reuse it in #4, e.g., `ut <- wind |> filter(...)`. Then run `freq()` on any variable in `ut` (e.g., `ut |> freq(Site.State)`) and read the **Total** row to get the number of turbines.

- 4) Using the subset you saved in #3 (e.g., `ut`), examine the frequency distribution of `Turbine.Capacity` (measured in kW).
 - a) How many turbines in this subset have a capacity of 2,300 kW?
 - b) What percentage of turbines in this subset have a capacity of 1,500 kW?
- 5) Often, we want to look at descriptive statistics of a quantitative variable. We can do this using the `descr()` function that is in the package `summarytools`. Examine the descriptive statistics for `Turbine.Total_Height` and answer the questions below. Read the note below the questions before you start. Include the units of measurement in your answers.
 - a) What is the mean?
 - b) What is the median?
 - c) What is the maximum?

i Note

The template for `descr()` is `descr(dataframe$variable)` (or `dataframe |> descr(variable)`, using the pipe). Unlike `freq()`, which counts cases in categories, `descr()` is used for **quantitative** (numeric) variables and produces a table of descriptive statistics, including the mean, median, minimum, maximum, and standard deviation. Find the row labeled with each statistic to answer the questions below, and check the codebook for the units (e.g., meters) that variable is measured in.

- 6) Next, we will learn how to use the function, `mutate()`. We use this function to create new variables (or “mutate” them) from existing ones. Let’s create a new variable called `Turbine_Space` that is the mean of two existing variables, `Turbine.Hub_Height` and `Turbine.Rotor_Diameter`. Then, examine the descriptive statistics of the new variable, `Turbine_Space`, and answer the following questions. Read the note below the questions before you start. Include units in your answers (both original variables are measured in meters).
- What is the mean of `Turbine_Space`?
 - What is the minimum value of `Turbine_Space`?
 - What is the maximum value of `Turbine_Space`?

i Note

The general template for `mutate()` is:

```
df <- df |>  
  mutate(new_variable_name = calculation_using_existing_variables)
```

Two things to keep in mind:

- **Save your changes.** Notice that the template starts with `df <-`. This reassigns the result back into your data frame. If you skip this and only run `df |> mutate(...)`, R will print the result to your Console but the new variable will not actually be saved — you won’t be able to use it in later steps. This is the same idea as when you first read your data into an object with `bway <- read_csv(...)` or `wind <- read_csv(...)`.
- **R calculates row by row automatically.** To find the mean of two variables *for each row* (i.e., for each wind turbine), you don’t need a special function — you can just add the two variables together and divide by 2, e.g., `(variable1 + variable2) / 2`. R automatically performs this calculation separately for every row in the dataset.

- 7) Sometimes, we may want to make continuous variables into categorical variables. For example, let’s say we are interested in the number of turbines in a project, `Project.Number_Turbines`, but we want to categorize projects as having a low or high number of turbines.

We refer to this process as **recoding variables**. For this exercise, let’s recode projects with 100 turbines or fewer into the **low** category and those that have more than 100 turbines into the **high** category.

Use the `mutate()` and `case_when()` functions to recode `Project.Number_Turbines` into a new variable, `Project.Size`, with only two categories, **low** (≤ 100) and **high** (> 100). Read the note below the questions before you start, then answer the following questions. Remember that each row is one **turbine**, not one project.

- How many turbines are in **low** projects (100 or fewer turbines)?
- How many turbines are in **high** projects (more than 100 turbines)?

i Note

The general template for combining `mutate()` and `case_when()` is:

```
df <- df |>
```

```
mutate(new_variable_name = case_when(  
  existing_variable <= value ~ "label1",  
  existing_variable > value ~ "label2"  
))
```

`case_when()` checks each row against your conditions, in order (using the same `==`, `>`, `<`, `>=`, and `<=` operators you've already used with `filter()`), and assigns the label after the `~` symbol to any row where that condition is `TRUE`. Just like with `mutate()` above, remember to reassign the result back to `wind` (i.e., start with `wind <- wind |> ...`) so that your new variable is saved and available for the questions below.

When you run `freq()` on your new variable, the counts tell you how many turbines belong to low or high projects. Note that `freq()` lists categories in alphabetical order, so `high` appears before `low`.

- 8) **BONUS (2 points):** In the previous step, you created a new variable called `Project.Size`. What is the mean `Project.Capacity` (measured in MW) for turbines in (a) low and (b) high projects?

Tip

One way to do this is to `filter()` for `low` and run `descr(Project.Capacity)`, then repeat for `high`, the same approach you've used throughout this assignment.

Optional: look into the function `group_by()`. Adding `group_by(variable)` to your pipe *before* `descr()` will calculate the statistics separately for each category of `variable`, all in a single step: `df |> group_by(variable) |> descr(other_variable)`.

Submission

Submit your R script (named `LA-#_FirstName-LastName.R`) to Canvas.

Your R script should:

- 1) Include commands and functions that are necessary to address all the questions in the assignment.
- 2) Contain comments that answer the questions in the assignment.
- 3) Run in its entirety without errors.

To ensure that your R script runs without errors, you should:

- Save your script.
- Navigate back to Your Workspace on Posit Cloud.
- Reopen your project.
- Run the entire script line-by-line without editing it to ensure there are no errors.

Important

These standards apply to all submissions in this course that require R scripts. You should follow these instructions for preparation, naming, and saving of your R script for all of your individual lab assignments.